Warden

Product Description · Spend guardrails and audit receipts for AI agents · Category: agent infrastructure / payments middleware (not an issuer, not a wallet, not fraud detection) · July 2026

> **ONE-LINER**
>
> Spend guardrails and audit receipts for AI agents that pay with virtual cards. Humans set the terms; Warden enforces them at the card network and proves why every dollar moved.

Problem

AI agents are beginning to spend real money, and two forces are colliding:

1. **Agents can't be trusted with open-ended access to funds.** Prompt injection is unsolved; in-the-wild payloads already try to coerce shopping agents into unauthorized purchases (Unit 42, 2026). Any spending check that runs inside the agent's own software can be routed around.
2. **The card networks now require accountability.** Visa's April 2026 rules (§4.1.24) and Mastercard's Verifiable Intent (March 2026) mandate intent capture and record retention for agentic payments.

Today nothing on the agent side enforces a human's budget non-bypassably, and nothing records *why* a charge happened at the only place the "why" exists: the agent's triggering instruction. That gap is the product.

Target user

Developers and small teams who run agents that spend money (shopping agents, procurement bots, ops automations). Single-tenant in v1: one operator, their agents, their money, their rules. The operator already has an AgentCard account; Warden layers on top.

What it is

A policy-and-audit proxy that sits between any agent and the AgentCard MCP server (the card-issuing rail). Three jobs:

1. **Enforce** human-set spending policy: budgets, merchant/category allow-block lists, per-merchant caps, velocity limits, approval thresholds.
2. **Contain** failures: every purchase gets its own single-use card ($1–$50, auto-cancels after one authorization), so the worst case for any failure is one small card.
3. **Explain** everything: every charge is bound to the agent's triggering intent and rendered as an immutable receipt, exportable in a format mapped to the Visa and Mastercard record duties.

Warden never holds funds, never stores card credentials, and never puts an ML model in the enforcement path.

How it works (user journey)

1. **Setup:** operator installs Warden, runs `warden auth` once to connect their AgentCard account, points their agent at Warden's MCP endpoint instead of AgentCard's. Ships with a sane default policy; day one requires nothing but a budget number.
2. **Delegation:** the agent calls `warden_start_task` with a stated intent ("restock office supplies") and gets a budget envelope.
3. **Purchase:** for each buy, the agent requests a card. Warden's deterministic policy engine evaluates the request (caps, lists, budget remaining, velocity, circuit state) and either mints a single-use card scoped to that amount, blocks with machine-readable reasons, or parks it for human approval.
4. **Oversight:** over-threshold purchases wait for a one-tap human approve/deny in the dashboard. A background reconciler ingests settled transactions every 30 seconds and writes receipts. If an agent goes runaway (e.g. 40 card requests in a minute), the circuit breaker closes all its open cards and halts its tasks automatically.
5. **Audit:** the operator sees every dollar per agent, per task, per intent, and can export the full receipt log as NDJSON.

Core capabilities

CAPABILITY	WHAT IT INCLUDES
Policy engine	Versioned, per-agent or global rules: allow/block merchants, category rules, per-card cap, per-task budget, per-merchant caps, velocity limits, approval threshold. Pure and deterministic.
Scoped card issuance	One single-use card per purchase, amount-capped at the network, auto-cancel after one authorization, TTL auto-close for unused cards.
Receipts	Append-only log joining transaction → card → task → intent → policy decision. Dashboard with full decision detail per receipt.
Human approvals	Threshold-triggered queue with approve/deny, expiry, optional webhook notification.
Circuit breaker	Hourly/daily velocity trip wires; auto close-all and halt with a dashboard "resume" action.
Multi-agent rollup	Per-agent policy, spend, receipts, blocks; org totals.
Audit export	NDJSON stream with fields mapped to Visa §4.1.24 record duties and the Mastercard Verifiable Intent schema.
Drop-in SDK	`withWarden(...)` one-line wrapper for any agent's payment flow.

THE ENFORCEMENT MODEL (KEY DESIGN POINT)

Software checks advise, the card enforces. Warden's policy decision compiles into the parameters of the card it mints. Dollar limits, single use, and expiry are enforced by the card network itself, so even an agent that bypasses Warden entirely cannot exceed one $50 authorization. Merchant and category rules are enforced deterministically at issuance (Warden refuses to mint), not at the network; the single-use card design is what bounds the damage if that layer is ever evaded. We say this plainly rather than overclaiming.

Differentiation

- **vs. Stripe Issuing / Lithic (issuer APIs):** they sell the card; Warden decides whether the card should exist and proves why it did. Issuer controls are card-level (MCC, amounts) and have no concept of agent, task, or intent. Building Warden's functionality on Stripe Issuing means building Warden.
- **vs. Visa Intelligent Commerce / Mastercard Agent Pay:** network-side protocols. The intent exists only on the agent side; those rules create the *demand* for agent-side capture, they don't supply it.
- **vs. Ramp Agent Cards:** corporate spend management for one company's employees-plus-agents; Warden is rail-level infrastructure for anyone's agents, rail-agnostic by design (the upstream is a swappable interface).
- **Whitespace:** agent-side intent capture + receipts. No one else can capture the prompt, because no one else is standing where the prompt is.

METRICS

- **North star:** receipts generated (charges bound to their triggering intent)
- **Secondary:** off-policy purchases blocked at the network
- Both counters posted publicly weekly during the build

SCOPE & CONSTRAINTS (V1)

- Sandbox (TEST) mode all season; live money only for the finale demo
- Single-tenant, single-machine, SQLite; no hosted service, no external SaaS
- No custody of funds; AgentCard owns issuance, wallets, KYC
- No ML anywhere in the enforcement path
- No checkout proxying: Warden controls card issuance and closure only
- Banned framing: "unhackable," "fraud detection," "AI security." Approved: guardrails, receipts, policy blocked an off-policy purchase

Known limitations (be upfront in reviews)

- Agent identity is self-declared (a name string), acceptable for a single operator, not cryptographic. Multi-tenant would need real agent credentials.
- Merchant rules are issuance-time, not network-time (see enforcement model).
- Reconciliation is polling-based (30s), so receipts and circuit-breaking are near-real-time, not instant.

Roadmap (6-week Ship Season, ships Aug 16, 2026)

WEEK	SHIP
1 · Jul 13–19	Receipts MVP: proxy live, every charge intent-bound, receipts dashboard
2 · Jul 20–26	Policy engine wired to issuance + policy editor UI
3 · Jul 27–Aug 2	Velocity caps + circuit breaker
4 · Aug 3–9	Human approvals + live prompt-injection demo (off-policy buy declines at the network)
5 · Aug 10–13	Multi-agent rollup + NDJSON audit export
6 · Aug 14–16	Drop-in SDK, landing page, pricing, waitlist, launch

Post-launch candidates: policy auto-drafting from observed spend (human approves the draft), category budgets and burn-down (currently a stretch goal), additional rails behind the swappable upstream interface.

Full technical spec (architecture, API contracts, 22 ordered build tasks): SPEC.md · Solo builder: Gagan Singh · Ship Season 2026