

Warden

Spend guardrails + receipts for AI agents, built on AgentCard · Ship Season 2026 · solo builder: Gagan Singh · 6-week roadmap, shipped Aug 16

PROBLEM

AI agents are starting to spend real money, and in-the-wild prompt-injection payloads already try to coerce agents into unauthorized purchases (Unit 42, 2026). Card networks now mandate intent capture and record retention for agentic payments (Visa Rules §4.1.24, Apr 2026; Mastercard Verifiable Intent, Mar 2026) - but nothing on the agent side enforces a human's budget non-bypassably or explains *why* a charge happened.

WHAT WARDEN DOES

A policy-and-audit layer wrapping the AgentCard MCP. Humans set budgets, merchant/category rules, and approval thresholds; Warden enforces them by minting single-use scoped cards (\$1–\$50, auto-cancel after one authorization) so declines happen at the card network - not in software an agent can route around. Every charge is bound to the agent's triggering intent and rendered as an explainable receipt. Warden never holds funds.

WEEK	SHIP	BUILD-IN-PUBLIC DEMO
1 · Jul 13–19	Receipts MVP. Wrap AgentCard MCP; every sandbox charge logged and bound to the agent's triggering intent; receipts dashboard.	"Every dollar my agent spent - with the why."
2 · Jul 20–26	Policy engine + scoped cards. Allow/block lists, per-merchant and category caps, task budgets; one single-use card per purchase, enforced at the network.	"My agent can only spend \$X at approved merchants."
3 · Jul 27–Aug 2	Velocity caps + circuit breaker. Hourly/daily limits; runaway agent auto-halted, all open cards closed.	"Agent tried 40 charges in a minute → frozen."
4 · Aug 3–9	Injection demo + human approvals. Live prompt-injection of my own shopping agent - the off-policy buy declines at the card network. Over-threshold buys need a human tap.	The viral moment (framed honestly: guardrails + receipts, not "unhackable").
5 · Aug 10–13	Multi-agent + audit export. Policy and rollup per agent; NDJSON receipts export mapped to Visa §4.1.24 record duties and the Mastercard Verifiable Intent schema.	"Which of my agents spent what - and prove it."
6 · Aug 14–16	Drop-in SDK + launch. One-line wrapper (<i>withWarden</i>) for any agent's payment flow;	Launch: receipts + blocks leaderboard totals.

landing page, pricing, waitlist; live demo at closing.

NORTH-STAR METRIC

Receipts generated - agent charges bound to their triggering intent.

- Secondary: off-policy purchases blocked at the network
- Weekly public posts of both counters

STACK & SCOPE GUARDRAILS

- TypeScript · MCP proxy · SQLite · Hono + React
- Sandbox (TEST) mode all season; live for finale
- Deterministic policy only - no ML in the enforcement path
- No custody of funds - AgentCard owns issuance/KYC

Full technical spec (architecture, API contracts, 22 ordered build tasks): SPEC.md · Market diligence: verified vs. Stripe Issuing, Lithic, Ramp Agent Cards, Visa Intelligent Commerce, Mastercard Agent Pay, x402 - whitespace is agent-side intent capture + receipts.