

Warden

Ship Season 2026 · Week 2 MVP Form — Prepared Answers · Gagan Singh · July 2026

Q1. Builder or team name

Gagan Singh (solo) — Warden

Q2. Email

gagan.2492@gmail.com

Q3. 60-second video showing and explaining your MVP (YouTube link)

[Paste YouTube link after recording]

Script, shot by shot (demo portion requires the Receipts MVP running in sandbox):

- 0:00–0:07 Hook** (screen: receipts dashboard). "My AI agent just spent money, and this receipt shows exactly why. I'm Gagan, building Warden."
- 0:07–0:17 Problem** (talking head or slide). "Agents are starting to spend real money. Prompt injection can trick them into buying things you never asked for, and Visa and Mastercard now require records of the human intent behind every agentic payment. Nothing on the agent side does that."
- 0:17–0:42 Demo** (screen recording). "Watch. My shopping agent starts a task: restock office supplies, budget 40 dollars. It asks Warden for a card. Warden mints a single-use card scoped to just this purchase; it dies after one charge. The agent buys. Thirty seconds later, here's the receipt: the charge, the card, the task, and the exact instruction that triggered it, in one record. Every dollar, with the why."
- 0:42–0:55 What's next** (slide). "This is week one of six. Coming: merchant and budget rules enforced at the card network, a circuit breaker for runaway agents, human approvals, and a live prompt-injection demo where the rogue purchase declines at the network."
- 0:55–1:00 Close.** "Warden. Your agents spend, you set the terms, every dollar explained. Follow along."

Q4. Repo / deploy URL

[GitHub repo link — push before submitting]

Q5. Tell us more about your idea and MVP: what it currently does and what is pending

Warden is spend guardrails plus audit receipts for AI agents that pay with virtual cards, built on the AgentCard MCP. The bet: agents can't be trusted with open-ended access to money (prompt injection is unsolved), and card networks now mandate intent capture for agentic payments (Visa §4.1.24, Mastercard Verifiable Intent). Warden is the agent-side layer that enforces a human's budget non-bypassably and records why every charge happened.

What the MVP does today (Receipts MVP): Warden sits as an MCP proxy between any agent and AgentCard. The agent declares a task with its triggering intent and a budget; each purchase mints a single-use card (\$1–\$50, auto-cancels after one authorization), so the blast radius of any failure is one small card. A background reconciler ingests every sandbox charge and binds it to the intent, producing an append-only receipt log with a dashboard: time, agent, merchant, amount, and the exact instruction behind the charge.

Pending over the next 5 weeks: the policy engine wired to issuance (merchant/category allow-block lists, per-merchant caps, task budgets) with a policy editor UI; velocity caps and a circuit breaker that auto-halts runaway agents; human approvals above a threshold; a live prompt-injection demo where the off-policy buy declines at the card network; multi-agent rollups; NDJSON audit export mapped to the Visa and Mastercard record duties; and a one-line SDK (withWarden).

The full product beyond this season: rail-agnostic (the issuer is a swappable interface, so Stripe Issuing or Lithic could sit underneath), multi-tenant with real cryptographic agent identity instead of self-declared names, hosted deployment, and policy auto-drafting where Warden observes spend and proposes rules the human approves with one tap. The MVP is deliberately single-operator, sandbox-mode, SQLite, and deterministic-only: the enforcement path will never contain an ML call.

Q6. Did anyone outside your team use or see the MVP?

Yes

Q7. What did you learn from your MVP being used or seen by other people?

Three pieces of pushback shaped my thinking. First: "humans don't set budgets on their own credit cards, they won't write policy rules." Partly right; the fix is that delegation always comes with terms (corporate cards, kids' cards, API spend caps), but policy authoring must be near-zero effort: sane defaults on day one, and later auto-drafted policies the human just approves or edits. Second: "Stripe already issues single-use cards." True, and it sharpened my positioning: the card is the commodity, the whitespace is agent-side intent capture and receipts; an issuer never sees the prompt that caused a charge, Warden does. Third, questions about enforcement forced me to be precise in all copy: dollar caps, single use, and expiry are enforced at the card network; merchant rules are enforced deterministically at issuance, and the single-use card bounds the damage if that layer is ever evaded.

Q8. Did your product roadmap change after the MVP exercise?

Yes

Q9. If it did, how did your roadmap change?

The 6-week ship order didn't change, but two things did. I added "policy auto-drafting" to the post-launch roadmap: Warden observes an agent's spend for a period, drafts a policy, and the human approves or edits it with one tap, which answers the real objection that nobody wants to hand-write merchant lists. And I re-weighted the messaging from guardrails-first to receipts-first: the intent-bound receipt is the thing no issuer or network can produce, so it leads every demo. One earlier change from the spec review: category budgets were downgraded to a stretch goal to protect the launch date.

Q10. Did you change your North-star Metric after the MVP exercise?

No

Still: receipts generated (agent charges bound to their triggering intent). Secondary: off-policy purchases blocked at the network.

Q11. If you did, tell us about your new North-star metric

Skip (no change).

Q12. What's your biggest blocker right now / what do you need from us?

Solo builder time is the main constraint. What would help most: intros to 3–5 people running agents that spend money today (shopping agents, procurement bots, ops automations) to act as design partners for the policy editor and receipts dashboard over the next two weeks.

Adjust to whatever is actually true when submitting.

Q13. Are you potentially looking to find a cofounder in the cohort?

[Your call: Yes / No]

Q14. If yes: what you're looking for in a cofounder

Draft, if answering yes: A distribution and GTM owner. I have the technical build covered end to end (spec, architecture, shipping weekly in public). The gap is someone who can own design partners, launch, pricing, and early sales conversations in fintech/payments while I ship, ideally someone who has sold developer tools or payments infrastructure before.

Q15. If yes: what you could offer a cofounder

Draft, if answering yes: A working product in a market with regulatory tailwind (Visa and Mastercard now mandate intent records for agentic payments), full technical ownership by a partner who ships fast, a public six-week build as proof of execution, and an equal seat in shaping the company from day one.